

# [v4,2/2] i386: Add support to get/set/migrate Intel Processor Trace feature

[diff \(/patch/881308/raw/\)](/patch/881308/raw/)[mbox \(/patch/881308/mbox/\)](/patch/881308/mbox/)[series \(/series/31863/mbox/\)](/series/31863/mbox/)

**Message ID** 1520182116-16485-2-git-send-email-luwei.kang@intel.com  
**State** New  
**Headers** show  
**Series** Untitled series #31863  
**Related** show

## Commit Message

Kang, Luwei (/project/qemu-devel/list/?submitter=69591)

March 4, 2018, 4:48 p.m.

**From:** Chao Peng <chao.p.peng@linux.intel.com>

Add Intel Processor Trace related definition. It also add corresponding part to kvm\_get/set\_msr and vmstate.

**Signed-off-by:** Chao Peng <chao.p.peng@linux.intel.com>

**Signed-off-by:** Luwei Kang <luwei.kang@intel.com>

---

```
target/i386/cpu.h      | 22 ++++++
target/i386/kvm.c      | 51 ++++++
target/i386/machine.c | 38 ++++++
3 files changed, 111 insertions(+)
```

## Patch

[diff \(/patch/881308/raw/\)](/patch/881308/raw/)[mbox \(/patch/881308/mbox/\)](/patch/881308/mbox/)[series \(/series/31863/mbox/\)](/series/31863/mbox/)

[patchwork.ozlabs.org/patch/881308/](https://patchwork.ozlabs.org/patch/881308/)

```

+                                     0x14, 1, R_EAX) & 0x7;
+
+     kvm_msr_entry_add(cpu, MSR_IA32_RTIT_CTL,
+                       env->msr_rtit_ctrl);
+     kvm_msr_entry_add(cpu, MSR_IA32_RTIT_STATUS,
+                       env->msr_rtit_status);
+     kvm_msr_entry_add(cpu, MSR_IA32_RTIT_OUTPUT_BASE,
+                       env->msr_rtit_output_base);
+     kvm_msr_entry_add(cpu, MSR_IA32_RTIT_OUTPUT_MASK,
+                       env->msr_rtit_output_mask);
+     kvm_msr_entry_add(cpu, MSR_IA32_RTIT_CR3_MATCH,
+                       env->msr_rtit_cr3_match);
+     for (i = 0; i < addr_num; i++) {
+         kvm_msr_entry_add(cpu, MSR_IA32_RTIT_ADDR0_A + i,
+                           env->msr_rtit_addrs[i]);
+     }
+ }
+
+     /* Note: MSR_IA32_FEATURE_CONTROL is written separately, see
+      *     kvm_put_msr_feature_control. */
@@ -2124,6 +2143,20 @@ static int kvm_get_msrs(X86CPU *cpu)
+ }
+
+ if (env->features[FEAT_7_0_EBX] & CPUID_7_0_EBX_INTEL_PT) {
+     int addr_num =
+         kvm_arch_get_supported_cpuid(kvm_state, 0x14, 1, R_EAX) & 0x7;
+
+     kvm_msr_entry_add(cpu, MSR_IA32_RTIT_CTL, 0);
+     kvm_msr_entry_add(cpu, MSR_IA32_RTIT_STATUS, 0);
+     kvm_msr_entry_add(cpu, MSR_IA32_RTIT_OUTPUT_BASE, 0);
+     kvm_msr_entry_add(cpu, MSR_IA32_RTIT_OUTPUT_MASK, 0);
+     kvm_msr_entry_add(cpu, MSR_IA32_RTIT_CR3_MATCH, 0);
+     for (i = 0; i < addr_num; i++) {
+         kvm_msr_entry_add(cpu, MSR_IA32_RTIT_ADDR0_A + i, 0);
+     }
+ }
+
+ ret = kvm_vcpu_ioctl(CPU(cpu), KVM_GET_MSRS, cpu->kvm_msr_buf);
+ if (ret < 0) {
+     return ret;
+ }
@@ -2364,6 +2397,24 @@ static int kvm_get_msrs(X86CPU *cpu)
+ case MSR_IA32_SPEC_CTRL:
+     env->spec_ctrl = msrs[i].data;
+     break;
+ case MSR_IA32_RTIT_CTL:
+     env->msr_rtit_ctrl = msrs[i].data;
+     break;
+ case MSR_IA32_RTIT_STATUS:
+     env->msr_rtit_status = msrs[i].data;
+     break;

```

```

+     case MSR_IA32_RTIT_OUTPUT_BASE:
+         env->msr_rtit_output_base = msrs[i].data;
+         break;
+     case MSR_IA32_RTIT_OUTPUT_MASK:
+         env->msr_rtit_output_mask = msrs[i].data;
+         break;
+     case MSR_IA32_RTIT_CR3_MATCH:
+         env->msr_rtit_cr3_match = msrs[i].data;
+         break;
+     case MSR_IA32_RTIT_ADDR0_A ... MSR_IA32_RTIT_ADDR3_B:
+         env->msr_rtit_addrs[index - MSR_IA32_RTIT_ADDR0_A] = msrs[i].data;
+         break;
+     }
+ }

```

diff --git a/target/i386/machine.c b/target/i386/machine.c

index 361c05a..c05fe6f 100644

--- a/target/i386/machine.c

+++ b/target/i386/machine.c

```

@@ -837,6 +837,43 @@ static const VMStateDescription vmstate_spec_ctrl = {
     }
 };

```

```

+static bool intel_pt_enable_needed(void *opaque)
+{
+    X86CPU *cpu = opaque;
+    CPUX86State *env = &cpu->env;
+    int i;
+
+    if (env->msr_rtit_ctrl || env->msr_rtit_status ||
+        env->msr_rtit_output_base || env->msr_rtit_output_mask ||
+        env->msr_rtit_cr3_match) {
+        return true;
+    }
+
+    for (i = 0; i < MAX_RTIT_ADDRS; i++) {
+        if (env->msr_rtit_addrs[i]) {
+            return true;
+        }
+    }
+
+    return false;
+}
+
+static const VMStateDescription vmstate_msr_intel_pt = {
+    .name = "cpu/intel_pt",
+    .version_id = 1,
+    .minimum_version_id = 1,
+    .needed = intel_pt_enable_needed,
+    .fields = (VMStateField[]) {
+        VMSTATE_UINT64(env.msr_rtit_ctrl, X86CPU),

```

```
+ VMSTATE_UINT64(env.msr_rtit_status, X86CPU),
+ VMSTATE_UINT64(env.msr_rtit_output_base, X86CPU),
+ VMSTATE_UINT64(env.msr_rtit_output_mask, X86CPU),
+ VMSTATE_UINT64(env.msr_rtit_cr3_match, X86CPU),
+ VMSTATE_UINT64_ARRAY(env.msr_rtit_addrs, X86CPU, MAX_RTIT_ADDRS),
+ VMSTATE_END_OF_LIST()
+ }
+};
+
+VMStateDescription vmstate_x86_cpu = {
+    .name = "cpu",
+    .version_id = 12,
@@ -957,6 +994,7 @@ VMStateDescription vmstate_x86_cpu = {
+    #endif
+    &vmstate_spec_ctrl,
+    &vmstate_mcg_ext_ctl,
+    &vmstate_msr_intel_pt,
+    NULL
+    }
+};
```

patchwork (<http://jk.ozlabs.org/projects/patchwork/>) patch tracking system | version 2.0.2.alpha-0 | [about patchwork \(/about/\)](#)